


Hello Minikube


Step 1: Create a minikube cluster

- a. Launch Docker Desktop and wait until it shows “Docker is running.”
- b. Open Windows PowerShell (or PowerShell as Administrator if needed).
- c. Run Minikube with:
 - i. `minikube start`



Step 2: Open the dashboard

- a. On a new terminal, run:
 - i. `minikube dashboard`
- b. Minikube will generate a link in the terminal.
- c. Click the link to open the Kubernetes Dashboard in your browser.



Step 3: Create a Deployment

- a. In your terminal, create a deployment by running:
 - i. `kubectl create deployment hello-node --image=registry.k8s.io/e2e-test-images/agnhost:2.53 -- /agnhost netexec --http-port=8080`

```
PS C:\Users\gab> kubectl create deployment hello-node --image=registry.k8s.io/e2e-test-images/agnhost:2.53 -- /agnhost netexec --http-port=8080
deployment.apps/hello-node created
```

- b. To check that your deployment was created successfully, run:
 - i. `kubectl get deployments`
 - c. It should display this output:

```
PS C:\Users\gab> kubectl get deployments
NAME        READY   UP-TO-DATE   AVAILABLE   AGE
hello-node  1/1     1           1           6m14s
```

- d. To see the pod created by your deployment, run:
 - i. `kubectl get pods`
 - e. It should display this output:

NAME	READY	STATUS	RESTARTS	AGE
hello-node-6c9b5f4b59-lh5tj	1/1	Running	0	8m16s

f. To view the cluster events, run:

i. `kubectl get events`

g. It should display this output:

LAST SEEN	TYPE	REASON	OBJECT	MESSAGE
8m21s	Normal	Scheduled	pod/hello-node-6c9b5f4b59-lh5tj	Successfully assigned default/hello-node-6c9b5f4b59-lh5tj to minikube
8m21s	Normal	Pulling	pod/hello-node-6c9b5f4b59-lh5tj	Pulling image "registry.k8s.io/e2e-test-images/agnhost:2.53"
8m1s	Normal	Pulled	pod/hello-node-6c9b5f4b59-lh5tj	Successfully pulled image "registry.k8s.io/e2e-test-images/agnhost:2.53" in n 20.125s (20.125s including waiting). Image size: 139374622 bytes
8m1s	Normal	Created	pod/hello-node-6c9b5f4b59-lh5tj	Created container: agnhost
8m	Normal	Started	pod/hello-node-6c9b5f4b59-lh5tj	Started container agnhost
8m22s	Normal	SuccessfulCreate	replicaset/hello-node-6c9b5f4b59	Created pod: hello-node-6c9b5f4b59-lh5tj
8m22s	Normal	ScalingReplicaSet	deployment/hello-node	Scaled up replica set hello-node-6c9b5f4b59 from 0 to 1
40m	Normal	Starting	node/minikube	Starting kubelet.
40m	Normal	NodeAllocatableEnforced	node/minikube	Updated Node Allocatable limit across pods
40m	Normal	NodeHasSufficientMemory	node/minikube	Node minikube status is now: NodeHasSufficientMemory
40m	Normal	NodeHasNoDiskPressure	node/minikube	Node minikube status is now: NodeHasNoDiskPressure
40m	Normal	NodeHasSufficientPID	node/minikube	Node minikube status is now: NodeHasSufficientPID
40m	Normal	RegisteredNode	node/minikube	Node minikube event: Registered Node minikube in Controller
40m	Normal	Starting	node/minikube	
30m	Normal	Starting	node/minikube	
30m	Normal	RegisteredNode	node/minikube	Node minikube event: Registered Node minikube in Controller

h. To view the kubectl configuration, run:

i. `kubectl config view`

j. It should display this output:

PS C:\Users\gab> kubectl config view	
apiVersion:	v1
clusters:	
- cluster:	
certificate-authority:	C:\Users\gab\.minikube\ca.crt
extensions:	
- extension:	
last-update:	Wed, 26 Nov 2025 10:56:21 +08
provider:	minikube.sigs.k8s.io
version:	v1.37.0
name:	cluster_info
server:	https://127.0.0.1:52589
name:	minikube
contexts:	
- context:	
cluster:	minikube
extensions:	
- extension:	
last-update:	Wed, 26 Nov 2025 10:56:21 +08
provider:	minikube.sigs.k8s.io
version:	v1.37.0
name:	context_info
namespace:	default
user:	minikube
name:	minikube
current-context:	minikube
kind:	Config
users:	
- name:	minikube
user:	
client-certificate:	C:\Users\gab\.minikube\profiles\minikube\client.crt
client-key:	C:\Users\gab\.minikube\profiles\minikube\client.key

j. To view application logs for a container in a pod, run this command:

i. `kubectl logs hello-node-6c9b5f4b59-lh5tj`

k. Make sure to replace 'hello-node-6c9b5f4b59-lh5tj' with the actual pod name from the 'kubectl get pods' command output.

```
PS C:\Users\gab> kubectl logs hello-node-6c9b5f4b59-lh5tj
I1126 03:18:47.122808      1 log.go:245] Started HTTP server on port 8080
I1126 03:18:47.124534      1 log.go:245] Started UDP server on port 8081
```

Step 4: Create a Service

- To expose the Pod to the public internet, run this command:
 - `kubectl expose deployment hello-node --type=LoadBalancer --port=8080`
 - type=LoadBalancer → Exposes the Pod outside the cluster (works with Minikube's internal LoadBalancer setup).
 - port=8080 → Maps to the container's port.

- It should display this output:

```
PS C:\Users\gab> kubectl expose deployment hello-node --type=LoadBalancer --port=8080
service/hello-node exposed
```

- To view the service created, run this command:

- `kubectl get services`

- It should display this output:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
hello-node	LoadBalancer	10.102.34.63	<pending>	8080:32390/TCP	9s
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	72m

- Lastly, run this command:

- `minikube service hello-node`

NAMESPACE	NAME	TARGET PORT	URL
default	hello-node	8080	http://192.168.49.2:32390

Starting tunnel for service hello-node./

NAMESPACE	NAME	TARGET PORT	URL
default	hello-node		http://127.0.0.1:61721

Starting tunnel for service hello-node.
Opening service default/hello-node in default browser...
! Because you are using a Docker driver on windows, the terminal needs to be open to run it.

- It would also open a browser window that serves as the response after running the command.



- It is important to note to only do this in a local Minikube cluster for testing. Do not expose `agnhost` with the `/shell` endpoint on a public or production cluster.

Step 5: Enable addons

- To view all the currently supported addons, run this command:
 - `minikube addons list`
- It should display this output

```
PS C:\Users\gab> minikube addons list
```

ADDON NAME	PROFILE	STATUS	MAINTAINER
ambassador	minikube	disabled	3rd party (Ambassador)
amd-gpu-device-plugin	minikube	disabled	3rd party (AMD)
auto-pause	minikube	disabled	minikube
cloud-spanner	minikube	disabled	Google
csi-hostpath-driver	minikube	disabled	Kubernetes
dashboard	minikube	enabled <input checked="" type="checkbox"/>	Kubernetes
default-storageclass	minikube	enabled <input checked="" type="checkbox"/>	Kubernetes
efk	minikube	disabled	3rd party (Elastic)
freshpod	minikube	disabled	Google
gcp-auth	minikube	disabled	Google
gvisor	minikube	disabled	minikube
headlamp	minikube	disabled	3rd party (kinvolk.io)
inaccel	minikube	disabled	3rd party (InAccel [info@inaccel.com])
ingress	minikube	disabled	Kubernetes
ingress-dns	minikube	disabled	minikube
inspektor-gadget	minikube	disabled	3rd party (inspektor-gadget.io)
istio	minikube	disabled	3rd party (Istio)
istio-provisioner	minikube	disabled	3rd party (Istio)
kong	minikube	disabled	3rd party (Kong HQ)
kubeflow	minikube	disabled	3rd party
kubetail	minikube	disabled	3rd party (kubetail.com)
kubevirt	minikube	disabled	3rd party (KubeVirt)
logviewer	minikube	disabled	3rd party (unknown)
metallb	minikube	disabled	3rd party (MetallLB)
metrics-server	minikube	disabled	Kubernetes
nvidia-device-plugin	minikube	disabled	3rd party (NVIDIA)
nvidia-driver-installer	minikube	disabled	3rd party (NVIDIA)
nvidia-gpu-device-plugin	minikube	disabled	3rd party (NVIDIA)
olm	minikube	disabled	3rd party (Operator Framework)
pod-security-policy	minikube	disabled	3rd party (unknown)
portainer	minikube	disabled	3rd party (Portainer.io)
registry	minikube	disabled	minikube
registry-aliases	minikube	disabled	3rd party (unknown)
registry-creds	minikube	disabled	3rd party (UPMC Enterprises)
storage-provisioner	minikube	enabled <input checked="" type="checkbox"/>	minikube
storage-provisioner-gluster	minikube	disabled	3rd party (Gluster)
storage-provisioner-rancher	minikube	disabled	3rd party (Rancher)
volcano	minikube	disabled	third-party (volcano)
volumesnapshots	minikube	disabled	Kubernetes
yakd	minikube	disabled	3rd party (marcnuri.com)

- To enable an addon, try running this command:
 - `minikube addons enable metrics-server`
- This starts the metrics-server Pod in the kube-system namespace.
- It allows you to view CPU and memory usage for your Pods using commands like:
 - `kubectl top pods`

```
PS C:\Users\gab> minikube addons enable metrics-server
💡 metrics-server is an addon maintained by Kubernetes. For any concerns contact minikube on GitHub.
You can view the list of minikube maintainers at: https://github.com/kubernetes/minikube/blob/master/OWNERS
  • Using image registry.k8s.io/metrics-server/metrics-server:v0.8.0
🌟 The 'metrics-server' addon is enabled
```

d. To View the Pod and Service you created by installing that addon:

i. `kubectl get pod,svc -n kube-system`

e. It should display this output:

PS C:\Users\gab> kubectl get pod,svc -n kube-system					
NAME	READY	STATUS	RESTARTS	AGE	
pod/coredns-66bc5c9577-p6dtg	1/1	Running	1 (85m ago)	95m	
pod/etcfd-minikube	1/1	Running	1 (85m ago)	95m	
pod/kube-apiserver-minikube	1/1	Running	1 (85m ago)	95m	
pod/kube-controller-manager-minikube	1/1	Running	1 (85m ago)	95m	
pod/kube-proxy-k48qf	1/1	Running	1 (85m ago)	95m	
pod/kube-scheduler-minikube	1/1	Running	1 (85m ago)	95m	
pod/metrics-server-85b7d694d7-2bdd4	1/1	Running	0	107s	
pod/storage-provisioner	1/1	Running	2 (85m ago)	95m	
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kube-dns	ClusterIP	10.96.0.10	<none>	53/UDP, 53/TCP, 9153/TCP	95m
service/metrics-server	ClusterIP	10.101.245.27	<none>	443/TCP	107s

f. To check the output from metrics-server, run this command:

i. `kubectl top pods`

g. It should display this output

PS C:\Users\gab> kubectl top pods		
NAME	CPU(cores)	MEMORY(bytes)
hello-node-6c9b5f4b59-lh5tj	1m	6Mi

h. If it displays 'error: Metrics API not available,' wait for a while, and try again.

i. To disable the metrics-server, run this command:

i. `minikube addons disable metrics-server`

j. It should display this output:

```
PS C:\Users\gab> minikube addons disable metrics-server
● "The 'metrics-server' addon is disabled
```

Get a Shell to a Running Container

Step 1: Getting a shell to a container

a. Create the pod by running this command:

i. `kubectl apply -f https://k8s.io/examples/application/shell-demo.yaml`

```
PS C:\Users\gab> kubectl apply -f https://k8s.io/examples/application/shell-demo.yaml
pod/shell-demo created
```

b. To verify if the container is running, run this command:

i. `kubectl get pod shell-demo`

```
PS C:\Users\gab> kubectl get pod shell-demo
NAME        READY   STATUS    RESTARTS   AGE
shell-demo  1/1     Running   0          27s
```

c. To get a shell into a running container, run this command:

i. `kubectl exec --stdin --tty shell-demo -- /bin/bash`

```
PS C:\Users\gab> kubectl exec --stdin --tty shell-demo -- /bin/bash
root@minikube:/# ls/
```

- d. The `--stdin` (or `-i`) keeps stdin open, and `--tty` (or `-t`) allocates a pseudo-terminal. The double dash `--` separates kubectl arguments from the command you want to run inside the container.
- e. Explore inside the container. Once in the shell, run commands like:
 - i. `ls /` - list the root directory
 - ii. `cat /proc/mounts` - view mounted filesystems
 - iii. `ps aux` - see running processes

```
root@minikube:/# ls /
bin  dev  docker-entrypoint.sh  home  lib64  mnt  proc  run  srv  tmp  var
boot docker-entrypoint.d  etc          lib   media  opt  root  sbin  sys  usr
```

Step 2: Writing the root page for nginx

- a. Try the nginx demo by creating a test HTML file by running this command:
 - i. `echo 'Hello shell demo' > /usr/share/nginx/html/index.html`
- b. To set a GET request to the nginx server, run this command:
 - i. `apt-get update`
 - `apt-get install curl`
 - `curl http://localhost/`
- c. It should display the "Hello shell demo"

```
root@minikube:/# apt-get update
apt-get install curl
curl http://localhost/
Get:1 http://deb.debian.org/debian trixie InRelease [140 kB]
Get:2 http://deb.debian.org/debian trixie-updates InRelease [47.3 kB]
Get:3 http://deb.debian.org/debian-security trixie-security InRelease [43.4 kB]
Get:4 http://deb.debian.org/debian trixie/main amd64 Packages [9670 kB]
Get:5 http://deb.debian.org/debian trixie-updates/main amd64 Packages [5412 B]
Get:6 http://deb.debian.org/debian-security trixie-security/main amd64 Packages [73.1 kB]
Fetched 9979 kB in 4s (2533 kB/s)
Reading package lists... Done
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
curl is already the newest version (8.14.1-2+deb13u2).
0 upgraded, 0 newly installed, 0 to remove and 0 not upgraded.
Hello shell demo
```

- d. Once done with shell, run the 'exit' command

```
root@minikube:/# exit
exit
```

Step 3: Running individual commands in a container

- a. In an ordinary command window, not shell, list the environment variables in the running container by running this command:
 - i. `kubectl exec shell-demo - env`

```
C:\Users\gab>kubectl exec shell-demo -- env
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
HOSTNAME=minikube
KUBERNETES_SERVICE_PORT_HTTPS=443
KUBERNETES_PORT=tcp://10.96.0.1:443
KUBERNETES_PORT_443_TCP=tcp://10.96.0.1:443
KUBERNETES_PORT_443_TCP_PROTO=tcp
KUBERNETES_PORT_443_TCP_PORT=443
KUBERNETES_PORT_443_TCP_ADDR=10.96.0.1
HELLO_NODE_PORT_8080_TCP_PORT=8080
HELLO_NODE_PORT=tcp://10.102.34.63:8080
HELLO_NODE_PORT_8080_TCP=tcp://10.102.34.63:8080
HELLO_NODE_PORT_8080_TCP_ADDR=10.102.34.63
KUBERNETES_SERVICE_HOST=10.96.0.1
HELLO_NODE_SERVICE_PORT=8080
KUBERNETES_SERVICE_PORT=443
HELLO_NODE_SERVICE_HOST=10.102.34.63
HELLO_NODE_PORT_8080_TCP_PROTO=tcp
NGINX_VERSION=1.29.3
NJS_VERSION=0.9.4
NJS_RELEASE=1~trixie
PKG_RELEASE=1~trixie
DYNPKG_RELEASE=1~trixie
HOME=/root
```

- b. Experiment with running other commands such as:

- i. `kubectl exec shell-demo -- ps aux`

```
C:\Users\gab>kubectl exec shell-demo -- ps aux
USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND
root      1  0.0  0.2 14764  8832 ?        Ss  04:11  0:00 nginx: master process nginx -g daemon off;
nginx     29  0.0  0.0 15224  3848 ?        S   04:11  0:00 nginx: worker process
nginx     30  0.0  0.0 15224  3592 ?        S   04:11  0:00 nginx: worker process
nginx     31  0.0  0.0 15224  3592 ?        S   04:11  0:00 nginx: worker process
nginx     32  0.0  0.0 15224  3464 ?        S   04:11  0:00 nginx: worker process
nginx     33  0.0  0.0 15224  3592 ?        S   04:11  0:00 nginx: worker process
nginx     34  0.0  0.0 15224  3464 ?        S   04:11  0:00 nginx: worker process
nginx     35  0.0  0.0 15224  3464 ?        S   04:11  0:00 nginx: worker process
nginx     36  0.0  0.0 15224  3464 ?        S   04:11  0:00 nginx: worker process
nginx     37  0.0  0.0 15224  3464 ?        S   04:11  0:00 nginx: worker process
nginx     38  0.0  0.0 15224  3464 ?        S   04:11  0:00 nginx: worker process
nginx     39  0.0  0.0 15224  3464 ?        S   04:11  0:00 nginx: worker process
nginx     40  0.0  0.0 15224  3464 ?        S   04:11  0:00 nginx: worker process
nginx     41  0.0  0.0 15224  3464 ?        S   04:11  0:00 nginx: worker process
nginx     42  0.0  0.0 15224  3592 ?        S   04:11  0:00 nginx: worker process
nginx     43  0.0  0.0 15224  3464 ?        S   04:11  0:00 nginx: worker process
nginx     44  0.0  0.0 15224  3464 ?        S   04:11  0:00 nginx: worker process
root     111  0.0  0.0 4332  3456 pts/0    Ss+  05:02  0:00 /bin/bash
root     237 25.0  0.0  6396  3584 ?        Rs   05:02  0:00 ps aux
```

ii. `kubectl exec shell-demo -- ls /`

```
C:\Users\gab>kubectl exec shell-demo -- ls /
bin
boot
dev
docker-entrypoint.d
docker-entrypoint.sh
etc
home
lib
lib64
media
mnt
opt
proc
root
run
sbin
srv
sys
tmp
usr
var
```

iii. `kubectl exec shell-demo -- cat /proc/1/mounts`

```
C:\Users\gab>kubectl exec shell-demo -- cat /proc/1/mounts
overlay / overlay rw,relatime,lowerdir=/var/lib/docker/overlay2/l/73LNMF57B5AVURIU5JXA65LG:/var/lib/docker/overlay2/l/63TQMVEEOHA2PMMIOPXK2E2X7:/var/lib/
docker/overlay2/l/WVXG0BIDVETL3T2QE7UHF33P4:/var/lib/docker/overlay2/l/033LA4UJJSFRXGUC3YMSXYLHAB3:/var/lib/docker/overlay2/l/2K5W4LDSVM6ZL4XHW6TUBWYJGP:/va
r/lib/docker/overlay2/l/GUUTESQDRSG7H4JRCJLHWJ51W:/var/lib/docker/overlay2/l/AFBGUSYZACGWHYCQGJBCPZ56G:/var/lib/docker/overlay2/l/NLHRNTW7Q7AEFVSXPQL3ZMI
4H,upperdir=/var/lib/docker/overlay2/3/c440b9509408221c9a123692fbeaaedea81da271c5e6735e3ce156c59ae733/work 0
proc /proc rw,nosuid,nodev,noexec,relatime 0 0
tmpfs /dev tmpfs rw,nosuid,size=65536k,mode=755 0 0
devpts /dev/pts devpts rw,nosuid,noexec,relatime,gid=5,mode=620,ptmxmode=666 0 0
sysfs /sys sysfs ro,nosuid,nodev,noexec,relatime 0 0
cgroup /sys/fs/cgroup cgroup2 ro,nosuid,nodev,noexec,relatime 0 0
mqueue /dev/mqueue mqueue rw,nosuid,nodev,noexec,relatime 0 0
shm /dev/shm tmpfs rw,nosuid,nodev,noexec,relatime,size=65536k 0 0
/dev/sda /dev/termination-log ext4 rw,relatime 0 0
/dev/sde /etc/resolv.conf ext4 rw,relatime 0 0
/dev/sde /etc/hostname ext4 rw,relatime 0 0
/dev/sde /etc/hosts ext4 rw,relatime 0 0
/dev/sde /usr/share/nginx/html ext4 rw,relatime 0 0
tmpfs /run/secrets/kubernetes.io/serviceaccount tmpfs ro,relatime,size=3858624k,noswap 0 0
proc /proc/bus proc ro,nosuid,nodev,noexec,relatime 0 0
proc /proc/fs proc ro,nosuid,nodev,noexec,relatime 0 0
proc /proc/irq proc ro,nosuid,nodev,noexec,relatime 0 0
proc /proc/sys proc ro,nosuid,nodev,noexec,relatime 0 0
proc /proc/sysrq-trigger proc ro,nosuid,nodev,noexec,relatime 0 0
tmpfs /proc/acpi tmpfs ro,relatime 0 0
tmpfs /proc/interrupts tmpfs rw,nosuid,size=65536k,mode=755 0 0
tmpfs /proc/kcore tmpfs rw,nosuid,size=65536k,mode=755 0 0
tmpfs /proc/keys tmpfs rw,nosuid,size=65536k,mode=755 0 0
tmpfs /proc/latency_stats tmpfs rw,nosuid,size=65536k,mode=755 0 0
tmpfs /proc/timer_list tmpfs rw,nosuid,size=65536k,mode=755 0 0
tmpfs /proc/scsi tmpfs ro,relatime 0 0
tmpfs /sys/firmware tmpfs ro,relatime 0 0
```

c. When a Pod contains multiple containers, the `--container` (or `-c`) flag must be provided with the `kubectl exec` command to specify which container to access.

- i. For example, consider a Pod named `my-pod` with two containers: `main-app` and `helper-app`.
- ii. The following command opens a shell inside the `main-app` container:
- iii. `kubectl exec -i -t my-pod --container main-app -- /bin/bash`

Deploying WordPress and MySQL with Persistent Volumes

Step 1: Create a working directory

- Create a new directory named wordpress-k8s and move into it:

- `mkdir wordpress-k8s`
- `cd wordpress-k8s`

```
Windows PowerShell
Copyright (C) Microsoft Corporation. All rights reserved.

Install the latest PowerShell for new features and improvements! https://aka.ms/PSWindows

PS C:\Users\gab> mkdir wordpress-k8s

Directory: C:\Users\gab

Mode                LastWriteTime         Length Name
----                <-----              ----- 
d----       27/11/2025   1:11 pm            wordpress-k8s

PS C:\Users\gab> cd wordpress-k8s
```

Step 2: Create a kustomization.yaml

- There are two ways to create the kustomization.yaml. First, you can create through notepad and paste this.

- ```
cat <<EOF >./kustomization.yaml
secretGenerator:
- name: mysql-pass
 literals:
 - password=YOUR_PASSWORD
EOF
```

- Another method to create this file, is through window power shell by pasting this line:

- ```
@"
secretGenerator:
- name: mysql-pass
  literals:
  - password=mypass123
resources:
- mysql-deployment.yaml
- wordpress-deployment.yaml
"@ | Out-File -FilePath kustomization.yaml -Encoding utf8
```

```

PS C:\Users\gab\wordpress-k8s> @"
>> secretGenerator:
>> - name: mysql-pass
>>   literals:
>>     - password=mypass123
>> resources:
>>   - mysql-deployment.yaml
>>   - wordpress-deployment.yaml
>> @" | Out-File -FilePath kustomization.yaml -Encoding utf8

```

Step 3: Add resource configs for MySQL and WordPress

- Download the MySQL deployment configuration file by running this command in window power shell:
 - Invoke-WebRequest <https://k8s.io/examples/application/wordpress/mysql-deployment.yaml> -OutFile mysql-deployment.yaml
- Download the WordPress configuration file by running this command in window power shell:
 - Invoke-WebRequest <https://k8s.io/examples/application/wordpress/wordpress-deployment.yaml> -OutFile wordpress-deployment.yaml

```

PS C:\Users\gab\wordpress-k8s> Invoke-WebRequest -Uri https://k8s.io/examples/application/wordpress/mysql-deployment.yaml -OutFile mysql-deployment.yaml
PS C:\Users\gab\wordpress-k8s> Invoke-WebRequest -Uri https://k8s.io/examples/application/wordpress/wordpress-deployment.yaml -OutFile wordpress-deployment.yaml

```

Step 4: Verify all 3 files

- To verify that all 3 files exist, run the 'dir' command, and it should display all 3 files.

```
PS C:\Users\gab\wordpress-k8s> dir
```

Directory: C:\Users\gab\wordpress-k8s

Mode	LastWriteTime	Length	Name
-a---	27/11/2025 5:45 pm	142	kustomization.yaml
-a---	27/11/2025 5:41 pm	1442	mysql-deployment.yaml
-a---	27/11/2025 5:41 pm	1341	wordpress-deployment.yaml

- To double check, view the directory itself in the file explorer.

Name	Date modified	Type	Size
mysql-deployment	27/11/2025 5:41 pm	Yaml Source File	2 KB
wordpress-deployment	27/11/2025 5:41 pm	Yaml Source File	2 KB
kustomization	27/11/2025 5:45 pm	Yaml Source File	1 KB

Step 5: Make sure Minikube is running

- a. Run the following command to check the Minikube status:
 - i. `minikube status`
- b. Look for the line that shows the host status:
 - i. If it displays ‘host: Running’, continue to Step 6.
 - ii. If it displays ‘host: Stopped’, or any error message, start Minikube with:
 - i. `minikube start`
- c. Wait for Minikube to finish starting (typically 1–2 minutes).
- d. Minikube must be running in order to create and manage the Kubernetes resources used for the WordPress deployment.

```
PS C:\Users\gab\wordpress-k8s> minikube status
minikube
type: Control Plane
host: Running
kubelet: Running
apiserver: Running
kubeconfig: Configured
```

Step 6: Deploy WordPress and MySQL, and Verify

- a. Apply the WordPress and MySQL configuration by running this command:
 - i. `kubectl apply -k ./`
- b. It should display this output:

```
PS C:\Users\gab\wordpress-k8s> kubectl apply -k ./ 
secret/mysql-pass-782846ffc9 created
service/wordpress created
Warning: spec.SessionAffinity is ignored for headless services
service/wordpress-mysql created
persistentvolumeclaim/mysql-pv-claim created
persistentvolumeclaim/wp-pv-claim created
deployment.apps/wordpress created
deployment.apps/wordpress-mysql created
```

- c. To verify that the secret was created, run this command:
 - i. `kubectl get secrets`
- d. It should display this output:

```
PS C:\Users\gab\wordpress-k8s> kubectl get secrets
NAME          TYPE      DATA   AGE
mysql-pass-782846ffc9  Opaque    1      10s
```

- e. To verify that PersistentVolumeClaims were provisioned, check the PVCs by running this command:
 - i. `kubectl get pvc`

```
PS C:\Users\gab\wordpress-k8s> kubectl get pvc
NAME           STATUS  VOLUME                                     CAPACITY  ACCESS MODES  STORAGECLASS  VOLUMEATTRIBUTESCLASS  AGE
mysql-pv-claim  Bound   pvc-a56bb915-98a0-45a8-9c92-7bae94570f08  20Gi     RWO          standard      <unset>                60s
wp-pv-claim    Bound   pvc-4afff2191-d5a2-4529-b641-2d97aaed7a47  20Gi     RWO          standard      <unset>                60s
```

- f. Dynamic provisioning may take several minutes before the PVCs reach the Bound state.
- g. To verify that the pods are running, check the pod status by running this command:
 - i. `kubectl get pods`
- h. Both wordpress must say "Running" under STATUS and "1/1" under READY. Hence, it should display this output:

```
PS C:\Users\gab\wordpress-k8s> kubectl get pods
NAME           READY   STATUS    RESTARTS   AGE
hello-node-6c9b5f4b59-lh5tj   1/1     Running   1 (9h ago)  30h
shell-demo      1/1     Running   0          5h43m
wordpress-59dfd8d54d-rpkh2   1/1     Running   0          6m8s
wordpress-mysql-5cc66b8b8-gn5gw 1/1     Running   0          6m8s
```



- i. Pod creation may take a few minutes before reaching the Running state.
- j. To verify that the WordPress Service is active, check the service by running this command:
 - i. `kubectl get services wordpress`
- k. It should display this output:

```
PS C:\Users\gab\wordpress-k8s> kubectl get services wordpress
NAME        TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)        AGE
wordpress   LoadBalancer  10.96.43.155  <pending>       80:30860/TCP  6m24s
```

- l. In Minikube, LoadBalancer Services do not receive an external IP. The EXTERNAL-IP field will remain in a pending state.
- m. To retrieve the WordPress access URL, use Minikube to obtain the Service URL by running this command:
 - i. `minikube service wordpress --url`
- n. It should display the url:

```
PS C:\Users\gab\wordpress-k8s> minikube service wordpress --url
http://127.0.0.1:56295
```

- o. Access the WordPress setup page by copying the URL from the previous command and opening it in a web browser. The WordPress setup screen should appear.



Step 7: Cleaning up

- a. To remove all resources created by the kustomization—including the Secret, Deployments, Services, and PersistentVolumeClaims—run:
 - i. `kubectl delete -k ./`
- b. This command deletes every object defined in the kustomization directory.

- c. This prevents unused Pods, Services, Secrets, and PersistentVolumeClaims from remaining in the cluster. Leaving these resources behind can consume system resources, create unnecessary storage usage, and interfere with future deployments or testing.